

STACKS, QUEUES, PRIORITY QUEUES ([slides](#))

Stack

- LIFO: last in first out
- abstract → push, pop
- `java.util.Stack` class → adds and removes from end of `ArrayList`
 - `push()` → add element to stack, $O(1)$
 - `pop()` → get and remove last element in, $O(1)$
- could write your own with `ListNodes`
 - can write `peek()` → return next value without removing

Queue

- FIFO: first in first out (fair)
- abstract → enqueue (add element to queue), dequeue (remove first in element)
- `java.util.Queue` interface
 - `add()` → at end, $O(1)$
 - `remove()` → from front, $O(1)$
 - can implement with `LinkedList`
- could write your own with `ListNodes`
 - can write `peek()` → same code as diy Stack

PriorityQueue

- keeps values in sorted order or by priority, smallest out first
 - objects must be `Comparable` or provide `Comparator`
- abstractly, is a sorted list

BINARY SEARCH TREE ([slides](#))

priority queue

- inefficient diy design
 - invariant → keep list sorted
 - remove() → always remove first and update myFirst, O(1)
 - add() → O(N) because have to search for insertion
 - OR keep list unsorted → peek/remove() O(N), add O(1)
- java API uses binary heap → balanced best of both worlds
 - peek() → O(1)
 - remove() → O(log(N))
 - add() → O(log(N))

binary heap at a high level

- sorted list of nodes → ordered binary tree of nodes
- heap property: every node is less than or equal to its successors
 - values are smaller at top, smallest is the root
- shape property: tree is full (each has 2 children) except rightmost positions on the last level

binary tree comparisons

- ArrayList → fast (access) not very dynamic (changes)
 - O(1) get but O(N) add/remove except at end
- LinkedList → dynamic, not very fast
 - O(N) get, O(1) add/remove (once you get there)
- binary search tree → fast and dynamic
 - O(log(N)) search and O(log(N)) add/remove, assuming tree is balanced
- TreeSet/Map
 - O(log(N)) add, contains, put, get (not amortized)

- stored in sorted order, can get range of values in sorted order efficiently
- HashSet/Map
 - O(1) add, contains, put, get (amortized)
 - unordered, cannot get range efficiently

binary tree nodes

- has value and right+left child nodes

terminology

- root: top node, has no parent, node you pass for the whole tree
- leaf: bottom nodes, have no children (both null)
- path: sequence of parent-child nodes
- subtree: nodes at and beneath
- depth: number of edges (lines) from root to node (can have different conventions)
- height: maximum depth

binary search tree if

- left subtree values are less than this nodes value, right subtree values are all greater
- enables efficient search like binary search

TREE RECURSION ([slides](#))

ways to recursively traverse a binary tree

- inOrder: left, root, right
- preOrder: root, left, right
- postOrder: left, right, root

ex: TreeCount

- base case \rightarrow tree == null, return 0
- recursive case \rightarrow $1 + \text{count}(t.\text{left}) + \text{count}(t.\text{right})$

complexity of tree traversal

- create recurrence relation and solve ($T(N) = a*T(g(N)) + f(N)$)
- tree traversal $\rightarrow O(n)$ regardless of balance
 - balanced $\rightarrow T(n/2) + O(1) + T(n/2) = O(n)$
 - $T(n/2)$ because same number left and right
 - unbalanced $\rightarrow T(0) + O(1) + T(n-1) = O(n)$
 - if every node has one child on one side
- search \rightarrow balanced $O(\log(n))$, unbalanced $O(n)$
 - balanced $\rightarrow T(n/2) + O(1) = O(\log(n))$
 - unbalanced $\rightarrow T(n-1) + O(1) = O(n)$
- a binary tree (a,b) is approximately balanced if
 - for every node rooting a subtree of size $n \geq a$, left and right subtrees have at most $b(n/2)$ nodes
- approximately balanced tree is good enough for $O(\log(N))$

GREEDY ALGORITHMS FOR DISCRETE OPTIMIZATION (slides)

optimization

- find solution that maximizes or minimizes an objective

greedily searching

- start with partial solution, take a step toward a complete solution in each iteration
- greedy principle: in each iteration, make the lowest cost or highest value step
- does not always guarantee the best overall solution (global optima)
- sometimes is optimal or not provably optimal but works in practice
- ex: machine learning → in a neural network, make a small change to best improve performance

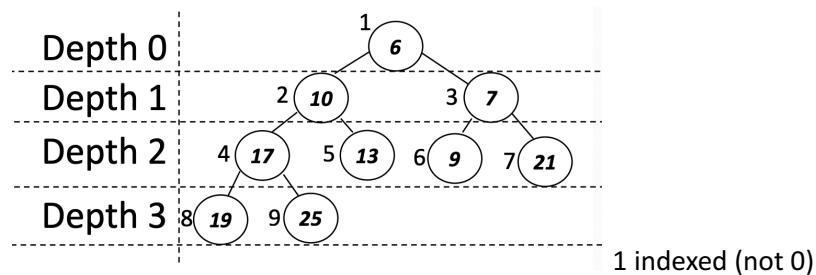
BINARY HEAPS ([slides](#))

binary heap

- java API PriorityQueue is implemented as a binary heap
- heap property: every node is less than or equal to its successors
 - values are smaller at top, smallest is the root
- shape property: tree is full (each has 2 children) except rightmost positions on the last level

using an array

- PriorityQueue actually implements with an array
 - minimizes storage (no nodes)
 - simpler to code (no explicit tree traversal)
 - faster (constant, not asymptotic) → children are located by index in array
- for a Heap
 - index positions into the tree level by level, left to right



- last node in heap is at the max index, min node is at index 1
- peek returns first value
- with 1 indexing, for node with index k
 - left child index = $2k$
 - right child index = $2k+1$
 - parent index = $k/2$
- heap is complete → complete binary tree has 2^d nodes at depth d

heap add

- add to first open position at last level of tree (end of array)
- swap with parent if heap property violated → stop when parent smaller or root reached

heap remove

- returns root value
- replace root with last node in heap
- while heap property violated (while node is larger than current node), swap with smaller child

BALANCED BINARY SEARCH TREES ([slides](#))

binary search tree invariant

- binary tree that for every node → left subtree values < node, right subtree values > node

diy TreeSet

- stores no duplicates in sorted order → prints the tree values in-order
- add() → wrapper for recursive helper insert()
 - if less, insert or recurse left; if greater, insert or reuse right; if equal, nothing (duplicate)
- contains() → wrapper for recursive search
 - base cases → no more tree to search, did not find target; found target
 - if target < node, search left; if target > node, search right
- runtime for add/contains on balanced tree
 - $T(N) = T(N/2) + O(1)$, $O(\log(N))$ → same as binary search
- runtime for add/contains on (perfectly) unbalanced tree
 - $T(N) = T(N-1) + O(1)$ → $O(N)$, search in linked list

runtime

- peek() → return first value of array, $O(1)$
- add(), remove() → traverse one root-leaf path, max is $\log N$
 - complete binary tree always has height $O(\log(N))$
- runtime of add/contains is proportional to height of the binary search tree

red-black tree

- criteria (for binary search tree)
 - root is black, red node can't have red children
 - all paths from a node to null nodes must have same number of black nodes
- TreeMap is a red-black tree based in NavigableMap implementation
- TreeSet is a NavigableSet implementation based on a TreeMap

- not all binary search trees can be colored as a red-black trees → if not approximately balanced
- properties guarantee that approximate balance (good enough for $O(\log N)$)
 - any red-black tree with N nodes has height $O(\log N)$
- functionality
 - contains/search → same as bst
 - add/insert → run regular bst add/insert, color new node red, recolor tree
 - sometimes won't work → need to rotate (change structure, change root)

GRAPHS, DFS, BFS ([slides](#), [slides](#))

graph

- data structure for representing connections among items, vertices connected by edges
- vertex (node) represents item, edge represents connection between 2 vertices
- undirected graph: edges go both ways (ex: webpage and links)
- directed graph: edges go one way (social media networks)
- simple graph: at most one (undirected) edge between nodes (or 2 directed)
- size of graph
 - N or $|V| \rightarrow$ number of edges
 - M or $|E| \rightarrow$ number of edges
 - in a simple graph, $M \leq N^2$
- path: sequence of unique vertices connected by edges (or edges with unique vertices)

recursive depth-first search (DFS) in grid graphs

- 2d grid is a graph with implicit structure
- maze is a grid graph (2d with not all edges present)
 - edge = no wall, no edge = wall
- depth first search for solving maze
 - always explore (recurse on) a new (unvisited) adjacent vertex if possible
 - if not, backtrack to the most recent vertex adjacent to an unvisited vertex and continue
 - base cases \rightarrow searching off the grid, already visited here (keep in array to avoid infinite)

```
176     if (!north[x][y]) {  
177         int d = solveDFS(x, y + 1, depth+1);  
178         if (d > 0) return d;  
179     }  
  
    !north[x][y] → no wall  
    above, can go that way.  
    y+1 → recurse  
    on node above  
    Tracking length  
    of path  
  
    If you found the center,  
    return the path length
```

- recursive case \rightarrow 3 more symmetric cases for other 3 directions
- N width x height nodes \rightarrow runtime $O(N)$

- each node will be recursed on ≤ 4 times (4 neighbors that can recurse on it)
- each recursive call is $O(1)$
- vs recursive tree traversal → tree traversal assume 2 adjacent nodes and no cycles

iterative graph search data structures

- have an adjacency list for the graph → `HashMap<Vertex, HashSet<Vertex>> aList`
- keep track of visited nodes → `Set<Vertex> visited`
- keep track of the previous node → `Map<Vertex, Vertex>`

iterative DFS

- stack to store unexplored nodes (LIFO)
- keep going on one path until it can't anymore
- iterative DFS loop structure
 - while `toExplore` not empty, explore from the recently discovered node (`current`)
 - look at all the neighbors and check if unvisited (not in `visited`)
 - note how we got to this neighbor (put in `previous`), note we've seen this (add `visited`), mark to explore (push to `toExplore`)
- DFS search tree → can find paths from one point to another by going backwards from the end point, the particular paths found in this graph
- DFS complexity (for N vertices, M edges)
 - pop each of N nodes at most once
 - loop over neighbors of each node exactly once, consider each edge twice (each relationship is recorded twice)
 - $N+2M \rightarrow O(N+M)$

iterative breadth-first search (BFS)

- queue to store unexplored nodes (FIFO)
- explore all your neighbors before visiting neighbors' neighbors
- iterative BFS loop structure
 - while `toExplore` not empty, explore from the closest discovered node (`current`)

- look at all the neighbors and check if unvisited (not in visited)
- note how we got to this neighbor (put in previous), note we've seen this (add visited), mark to explore (push to toExplore)
- BFS guarantees the shortest path, DFS does not → BFS does not explore all paths, only shortest

WEIGHTED GRAPHS, DIJKSTRA ([slides](#))

weighted graphs

- each edge has an associated weight representing cost, distance, etc

Dijkstra data structures

- have an adjacency list for the graph → stores neighbors for each vertex (aList)
 - `HashMap<Vertex, HashSet<Vertex>> aList` → $O(1)$ to get neighbors with good hashCode
- ~~keep track of visited nodes in a set (visited)~~
- keep track of the previous node → how did i get here (previous)
- priority queue to store unexplored nodes (toExplore)
- hashmap of distance of the shortest path so far, to a given node

Dijkstra's algorithm

- search loop structure
 - while `toExplore` not empty, explore from the closest unexplored node (`current`)
 - get weight on current to neighbor edge
 - for each neighbor, if no path found yet (`!distance.containsKey(n)`) or the path through `current` is shorter (`distance.get(n) > distance.get(curr) + weight`)
 - put this path (`distance.get(curr) + weight`) for neighbor in `distance`
 - note how we got to this neighbor (put in `previous`), note we've seen this (add `visited`), mark to explore (add to `toExplore`)
 - might add the same node to the priority queue multiple times
 - adding neighbor to `toExplore` usually updates the priority of neighbor, not add
 - ^ use `add` because most standard library binary heaps don't have efficient update
 - not guaranteed to be correct because greedy
 - runtime complexity
 - consider each node (N) once, each edge (M) twice, $\log(N)$ for each → $O((N+M)\log(N))$
 - heap duplicates

- while loop may loop more than N times
- in graphs with constant degree (constant number of neighbors), will still just be $O(N)$, but maybe not N
- worst case provable $O((N+M)\log(N))$ needs an efficient priority update

MINIMUM SPANNING TREE (MST), DISJOINT SETS, UNION FIND ([slides](#), [slides](#))

minimum spanning tree (MST) problem

- given N nodes and M edges, each with a weight/cost
- find a set of edges that connect all the nodes with minimum total cost → tree

Prim's algorithm (greedy)

- initialize → choose an arbitrary vertex
- partial solution → MST connecting subset of vertices
- greedy step → choose cheapest edge that connects a new vertex to the partial solution
- algorithm greedily grows by choosing closest unconnected vertex

Kruskal's algorithm (greedy)

- initialize → all nodes in disjoint sets
- partial solution → forest of spanning trees in disjoint sets
- greedy step → choose cheapest edge that connects 2 disjoint sets/trees
- algorithm greedily grows by cheapest edge that connects disjoint sets/trees
- pseudocode (input of N nodes, M edges, M edge weights)
 - let MST to an empty set
 - let S be a collection of N disjoint sets, one per node
 - while S has more than 1 set:
 - let (u, v) be the minimum cost remaining edge
 - find which sets u and v are in; if not equal:
 - union the sets
 - add (u, v) to MST
 - return MST
- runtime → $O(M(\log(M)+C))$
 - while loop → worst case, over all M edges

- removing from binary heap within loop $\rightarrow O(\log(M))$
- C is time for union/find